

Getting Results with R

First import the mockdata:

```
>library(readr) # this loads the library that knows how to read files
>mockdata <- read_csv("http://github.org/
  jon-may/GettingResultsinR/raw/master/mockdata.csv")
```

And tell R to use this data as the default 'data frame' so you do not need to specify it in every command:

```
>attach(mockdata)
```

counting cases

First, count how many males and females were tested, using `table`:

```
>table(Sex)
Sex
  F  M
20 10
```

Report it like this, in the *Participants* section of the *Method*:

“We recruited 30 participants (20 females, 10 males)...”

Descriptive statistics

Find out how old the participants were:


```
>library(psych) # a library that contains handy functions for psychology
>describe(Age)
  vars  n mean  sd median trimmed  mad min max range skew kurtosis  se
x1     1  30 20.9 3.03    20   20.38 1.48  18  29    11  1.26    0.59 0.55
```

“... aged between 18 and 29 (M=20.9, SD=3.0).”

As long as the skew and kurtosis are less than ± 2 the data can be treated as normally distributed, but that skew value is a bit high. Draw a histogram to check:

```
>hist(Age) # look in Plot window!
```

This shows why the skew value is so high, but luckily there is no need to do any parametric statistics on this demographic variable.



Chi squared test of association – χ^2

`table` can make two-way tables too. For example, how many of each sex were tested by each experimenter:

```
> table(Sex,Experimenter)
  Experimenter
Sex  JM  LS  MK
  F   6   6   8
  M   4   4   2
```

MK didn't seem to test as many men as the others – was this systematic, or just random variation?

Use a Chi Squared test to see if there is an association between Sex and Experimenter, and do so by sending it the results of this `table` command:

```
> chisq.test(table(Sex,Experimenter))
```

Pearson's Chi-squared test

```
data: table(Sex, Experimenter)
X-squared = 1.2, df = 2, p-value = 0.5488
```

A warning message is also produced, because some of the 'expected' cell sizes are less than 5, but there is no association. There is no need to report this, but this is how a χ^2 test is reported:

"There was no difference in the proportion of men and women tested by each experimenter $\chi^2(2, N=30)=1.20, p=.549$."

Correlation

The participants have been measured three times, at a baseline session (Value1), after training (Value2) and at a follow-up session (Value3). These can be correlated a pair at a time like this:

```
> cor(Value1,Value2)
[1] 0.3078597
```

With more than two variables, a correlation matrix can be produced by using cor on a data frame or part of a data frame, e.g.,

```
> cor(mockdata[,6:8])
      Value1      Value2      Value3
Value1 1.0000000  0.30785968  0.17061310
Value2 0.3078597  1.00000000 -0.02317446
Value3 0.1706131 -0.02317446  1.00000000
```

A subset of a dataframe

A subset of a data frame is referred to by putting row and column indices in square brackets after the name of the data frame. Here `mockdata[,6:8]` selects all rows (as there is nothing before the comma) and columns 6 to 8 (inclusive). If the indices are not known, they can be found by combining the `which` and `colnames` functions, e.g.:

```
> which(colnames(mockdata)=="Value1")
[1] 6
```

These can be used directly in the `cor` command:

```
> cor(mockdata[,
      which(colnames(mockdata)=="Value1"):
      which(colnames(mockdata)=="Value3")])
```

If the variables that are to be correlated are not adjacent, then this will not work, so a new data frame needs to be built that contains them:

```
> myvals<-mockdata[,c(1,6:8)]
> cor(myvals)
      Subject      Value1      Value2      Value3
Subject 1.0000000  0.2115570  0.26509522  0.52041127
Value1 0.2115570  1.0000000  0.30785968  0.17061310
Value2 0.2650952  0.3078597  1.00000000 -0.02317446
Value3 0.5204113  0.1706131 -0.02317446  1.00000000
```

The `c()` command joins together the values is encloses into a list which the `[,]` then uses to select the columns in the order specified. Here it selects the first column, then columns 6 to 8.

A limitation of `cor` is that it just gives the *r* value. The `psych` library contains a function `corr.test` (note that there are two *rs* in its name; there is another function `cor.test`).

```
> corr.test(myvals)
Call:corr.test(x = myvals)
Correlation matrix
      Subject Value1 Value2 Value3
Subject  1.00  0.21  0.27  0.52
Value1   0.21  1.00  0.31  0.17
Value2   0.27  0.31  1.00 -0.02
Value3   0.52  0.17 -0.02  1.00
Sample Size
[1] 30
Probability values (Entries above the diagonal are adjusted for multiple tests.)
      Subject Value1 Value2 Value3
Subject  0.00  0.79  0.63  0.02
Value1   0.26  0.00  0.49  0.79
Value2   0.16  0.10  0.00  0.90
Value3   0.00  0.37  0.90  0.00
```

To see confidence intervals of the correlations, print with the `short=FALSE` option
 There is nothing much to report here, but this is how it might be reported:

“The three scores did not correlate with each other (Baseline, After testing: $r(30) = .31$, $p=.49$; After testing, Follow-up: $r(30)=-.02$, $p=.90$; Baseline, Follow-up: $r(30)=.17$, $p=.79$; p values adjusted for multiple comparisons).”


There is a presumably spurious correlation between the Subject and Value3:

“There was a correlation between Subject Code and participant’s score at follow-up $r(30)=.52$, $p=.02$ ”

Scatterplots and regression

To understand a correlation it is always a good idea to draw a scatterplot. The `plot(x,y)` command draws the scatterplot, and then the `abline()` command adds the regression line using the result of a simple linear regression command `lm(y~x)`. Note that the order of x and y are different in these two commands!

```
> plot(Subject,Value3)
> abline(lm(Value3 ~ Subject))
```



The full output from the `lm()` command looks like this:

```
> summary(lm(Value3 ~ Subject))

Call:
lm(formula = Value3 ~ Subject)

Residuals:
    Min       1Q   Median       3Q      Max
-53.493 -27.454   1.567  24.302  61.759

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  93.1149    11.7032   7.956 1.15e-08 ***
Subject       2.1259     0.6592   3.225  0.0032 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 31.25 on 28 degrees of freedom
Multiple R-squared:  0.2708,    Adjusted R-squared:  0.2448
F-statistic: 10.4 on 1 and 28 DF,  p-value: 0.003197
```

To check that this is based on the same data as the correlation, the square root of the R^2 value should be .52

```
> sqrt(.2708)
[1] 0.5203845
```

The full regression equation statistics could be reported like this:

“A simple linear regression model showed that follow-up scores were predicted by Subject code, where Follow-up = $93.11 + \text{Subject} \times 2.13$, $R^2=.27$, $F(1,28)=10.4$, $p=.003$ ”.

Computing and Recoding variables

Existing variables may need to be combined for further analysis:

```
> mockdata$sumOfValues <- Value1 + Value2 + Value3
> mockdata$meanValue <- mockdata$sumOfValues / 3
```

Variables may need to be recoded from raw data into categories:

```
> mockdata$ValueGroup[mockdata$sumOfValues>0] <- "Low"
> mockdata$ValueGroup[mockdata$sumOfValues>250] <- "Medium"
> mockdata$ValueGroup[mockdata$sumOfValues>350] <- "High"
> describeBy(mockdata$sumOfValues, group=mockdata$ValueGroup, mat=TRUE, digits=1)
```

	item	group	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
X11	1	High	1	8	395.9	36.6	394	395.9	51.1	358	444	86	0.2	-1.9	12.9
X12	2	Low	1	5	206.6	32.1	225	206.6	17.8	167	237	70	-0.3	-2.2	14.4
X13	3	Medium	1	17	312.0	25.3	319	313.1	26.7	260	347	87	-0.5	-0.9	6.1

“Participants were divided into groups based on their total score over the three sessions. The Low group (N=5) ranged from 167 to 237 (M=207, SD=32), Medium (N=17) from 260 to 247 (M=312, SD=25), and High (N=8) from 358 to 444 (M=396, SD=37).”

Note that in these commands the data frame `mockdata` has been specified explicitly, so that the computed variables are added to it. Even if `mockdata` has previously been made the default data frame with an `attach()` command, the new columns can not yet be used without the data frame prefix. To make them available first `detach()` and then `attach()` again:

```
> table(ValueGroup)
Error in table(ValueGroup) : object 'ValueGroup' not found
> detach(mockdata)
> attach(mockdata)
> table(ValueGroup)
ValueGroup
  High    Low Medium
    8     5    17
```

If you no longer need variables, they can be removed.

```
> mockdata <- mockdata[, 1:8]
```

This will retain all rows, but only the first eight columns, deleting the three just computed.

Comparing two means from one group with a paired t test

When people have been measured twice on a variable, so they have a pair of measures, the means can be compared using a paired t test.

```
> t.test(Value1, Value2, paired=TRUE, var.equal = TRUE)

Paired t-test

data: Value1 and Value2
t = -3.3941, df = 29, p-value = 0.002011
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -39.530416  -9.802918
sample estimates:
mean of the differences
      -24.66667
> describe(Value1)
  vars  n  mean    sd median trimmed  mad min max range skew kurtosis   se
X1    1 30 83.03 34.17   76.5   81.42 37.81  29 163   134 0.43    -0.6 6.24
> describe(Value2)
  vars  n  mean    sd median trimmed  mad min max range skew kurtosis   se
X1    1 30 107.7 33.49   101  107.33  43  43 165   122 0.06    -1.06 6.11
```

“Scores increased from 83.0 (34.2) at baseline to 107.7 (33.5) after training, $t(29)=3.39, p=.002$.”

The minus sign before the t value can always be ignored as it simply depends on the order the variables are listed.

Comparing two groups' means with a two sample t test

Where two groups have been measured on the same variable, their means can be compared using a two sample t test. First the data needs to be split into the two groups. This can be done by creating a data frame for each group, which is useful if more commands might need to be run on each group:

```
> control.df <- mockdata[Group=="control",]
> intervention.df <- mockdata[Group=="intervention",]
```

These commands will select all of the rows of mockdata where Group has the specified value, and copy them into the new data frames. Note that two equals signs are needed, and that there is nothing after the comma so that all columns (variables) are selected.

Having created the two separate sets of data, the t.test command can be used to refer to them explicitly using the dollar notation, which takes the format `dataframe$variable`

```
> t.test(control.df$Value1,intervention.df$Value1,var.equal = TRUE)
```

Two Sample t-test

```
data: control.df$Value1 and intervention.df$Value1
t = -0.16284, df = 28, p-value = 0.8718
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -28.06455  23.93121
sample estimates:
mean of x mean of y
 82.00000  84.06667
```

This can also all be done within the t.test command if there is no need to retain the two groups' data for later use:

```
> t.test(mockdata[Group=="control",]$Value1,
         mockdata[Group=="intervention",]$Value1,
         var.equal = TRUE)
```

Here, creating the two subsets of data is helpful because it allows the group descriptives to be obtained easily:

```
> describe(control.df$Value1)
  vars  n mean    sd median trimmed  mad min max range skew kurtosis  se
x1    1 15  82 38.97    73  80.92 54.86  29 149  120 0.16   -1.49 10.06
> describe(intervention.df$Value1)
  vars  n mean    sd median trimmed  mad min max range skew kurtosis  se
x1    1 15 84.07 29.96    78    81 26.69  45 163  118 1.01    0.68 7.74
```

although this can also be done using this command:

```
> describeBy(Value1, group=Group)
```

“The control group (M=82.0, SD=39.0) and the intervention group (M=84.1, SD=30.0) did not differ at baseline $t(28)=0.16$, $p=.872$ ”

Oneway between-subjects Analysis of Variance

When there are more than two groups or levels of any factor, ANOVA is needed, not multiple t tests.

Test whether the three experimenters recruited people who were different at baseline.

```
> oneway.test(Value1 ~ Experimenter, var.equal = TRUE)
```

One-way analysis of means

```
data: mockdata$Value1 and mockdata$Experimenter
F = 3.447, num df = 2, denom df = 27, p-value = 0.04643
```

In the brackets, the formula `Value1 ~ Experimenter` makes `Value1` the IV and `Experimenter` the DV. The `var.equal=TRUE` bit tells the analysis to assume that the SDs for each `Experimenter` are equivalent. There is a significant effect, reported like this:

“The participants recruited by each `Experimenter` did not all have the same means at baseline $F(2,27)=3.45$, $p=.046$.”

```
> describeBy(mockdata$Value1, group=mockdata$Experimenter)
group: JM
  vars  n mean    sd median trimmed   mad min max range skew kurtosis   se
X1    1 10 64.2 27.5   60.5   63.12 24.46  29 108   79 0.33   -1.37 8.7
-----
group: LS
  vars  n mean    sd median trimmed   mad min max range skew kurtosis   se
X1    1 10 83.6 27.39    85   85.5 37.06  34 118   84 -0.3   -1.33 8.66
-----
group: MK
  vars  n mean    sd median trimmed   mad min max range skew kurtosis   se
X1    1 10 101.3 38.61   92.5   99.38 47.44  55 163  108 0.27   -1.67 12.21
> with(mockdata, pairwise.t.test(Value1, Experimenter, p.adj="bonferroni", paired=F))
```

MK is clearly higher than JM, but what about LS? Post hoc t tests are needed to compare them, using Bonferroni correction because we are making three tests at once:

```
> pairwise.t.test(Value1, Experimenter, p.adj="bonferroni", paired=F)
```

Pairwise comparisons using t tests with pooled SD

data: Value1 and Experimenter

```
      JM      LS
LS 0.544 -
MK 0.042 0.664
```

P value adjustment method: bonferroni

This handy command just gives the *p* values, not the *t* values. It could be reported like this, using the descriptives too:

“Two-tailed t tests with bonferroni correction indicated that MK’s participants (M=101.3, SD=38.6) scored higher than JM’s (M=64.2, SD=27.5) *p*=.042, and that LS’s participants (M=83.6, SD=27.4) did not differ to MK’s *p*=.664 or JM’s *p*=.544.”

aov - A more general ANOVA command

Another, more general way to do ANOVA is to use this function:

```
> m<-aov(Value1 ~ Experimenter)
> summary(m)
              Df Sum Sq Mean Sq F value Pr(>F)
Experimenter  2   6887    3443   3.447 0.0464 *
Residuals    27  26972     999
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The *F* and *p* values are the same as in the `oneway` output, but here the results are being stored in an object called `m` instead of printing it right away because the output from `aov` is more complicated, and putting it into `m` allows other commands to be run as well, such as an effect size statistic and post hoc tests (the label ‘*m*’ is just a name; it could be called anything).

eta squared (η^2):

```
> library(lsr) # this library includes the etaSquared command
> etaSquared(m)
              eta.sq eta.sq.part
Experimenter 0.2033986   0.2033986
```

There are two values here: *etasq* (which is the variance explained by the effect) and the partial eta squared (which supposedly adjusts for the inclusion of multiple effects and

interactions in the model). As there is only one effect in this simple oneway anova, they are identical, and the write up would be:

“The participants recruited by each Experimenter did not all have the same means at baseline $F(2,27)=3.45$, $p=.046$, $\eta^2=.20$.”

Post hoc tests

You can also ask for Tukey’s Honestly Significant Difference posthoc test if you need to make unplanned comparisons between more than two levels of a between-subject factor:

```
> TukeyHSD(m)
  Tukey multiple comparisons of means
    95% family-wise confidence level

Fit: aov(formula = Value1 ~ Experimenter)

$Experimenter
      diff      lwr      upr      p adj
LS-JM 19.4 -15.64614  54.44614 0.3690517
MK-JM 37.1   2.05386  72.14614 0.0363937
MK-LS 17.7 -17.34614  52.74614 0.4337410
```

As with the Bonferroni t tests, MK and JM’s groups differ. Write this up like this:

“Tukey’s HSD tests conducted post hic indicated that MK’s participants ($M=101.3$, $SD=38.6$) scored higher than JM’s ($M=64.2$, $SD=27.5$) $p=.036$, and that LS’s participants ($M=83.6$, $SD=27.4$) did not differ to MK’s $p=.434$ or JM’s $p=.369$.”

These p values are lower than those from the pairwise t tests, because Tukey’s HSD is a more sensitive test.

Two-between ANOVA

The formula in `aov` can be extended to add in multiple between-subject factors, for example, testing Experimenter and Group at the same time, using the following ‘star’ notation to request the main effects and all interactions:

```
> m<-aov(Value1 ~ Experimenter * Group)
> summary(m)

              Df Sum Sq Mean Sq F value Pr(>F)
Experimenter    2   6887    3443   3.144 0.0613 .
Group            1     32      32   0.029 0.8656
Experimenter:Group 2     652     326   0.298 0.7451
Residuals      24  26288    1095

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

> etaSquared(m)

              eta.sq eta.sq.part
Experimenter    0.2033986074 0.207595400
Group            0.0009460812 0.001217089
Experimenter:Group 0.0192701293 0.024219193
```

“At baseline, there were no differences between the experimenter’s groups $F(2,24)=3.14$, $p=.061$, $\eta_p^2=.21$ or between the conditions $F(1,24)=0.03$, $p=.866$, $\eta_p^2=.00$, and there was no interaction of experimenter and condition $F(2,24)=0.29$, $p=.745$, $\eta_p^2=.02$ ”

The same approach works for more than two between-subject factors, but repeated measures designs with within-subject factors requires more work. The data needs to be reshaped.

Reorganising data (Wide to Long)

In mockdata, people were measured three times: Value1 at baseline, Value2 at the end of training, and Value3 in a follow-up session some weeks later. The data file includes all three measurements for each participant on a single line, as is convenient for data entry. This is called WIDE format.

To see the first few rows, use the `head` command:

```
> head(mockdata)
# A tibble: 6 × 8
  Subject Experimenter Age Sex Group Value1 Value2 Value3
  <dbl>      <chr> <dbl> <chr> <chr> <dbl> <dbl> <dbl>
1         3          JM   18    M control    32     99     46
2         6          LS   19    F control    34    120     71
3         8          LS   20    M control    60     91     76
4        14          MK   18    F control   149    131     78
5        13          MK   18    F control    73     80     84
6        11          MK   20    F control   130    145     85
```

This is nice and compact, but many analyses in R require LONG format, where a single column contains all of the measurements, and another column says which time they are from.

This is how to turn a wide data set into a long one:

```
> library(reshape2) # contains the reshaping commands
> longdata<-melt(data=mockdata, measure.vars = c("Value1","Value2","Value3"))
```

This creates a new dataset called `longdata`. The `c(...)` lists the three variables that are all the same measure taken at different times. This is how the new file looks:

```
> head(longdata)
  Subject Experimenter Age Sex Group variable value
1         3          JM   18    M control  Value1    32
2         6          LS   19    F control  Value1    34
3         8          LS   20    M control  Value1    60
4        14          MK   18    F control  Value1   149
5        13          MK   18    F control  Value1    73
6        11          MK   20    F control  Value1   130
```

The level names are now in the column `variable`, and the actual values are in the column `value`

Finally, make this the new default data frame:

```
> detach(mockdata)
> attach(longdata)
```

Renaming levels in a factor and creating factors from variables

The variable names that are now in `variable` are not very meaningful. They can be changed to more meaningful labels, but as `variable` is a factor, the `levels()` command needs to be used:

```

> levels(variable)[levels(variable)=="Value1"]<-"Baseline"
> levels(variable)[levels(variable)=="Value2"]<-"After Testing"
> levels(variable)[levels(variable)=="Value3"]<-"Follow-up"
> head(longdata)
  Subject Experimenter Age Sex   Group variable value
1       1             JM  19   F control Baseline  108
2       2             JM  18   F control Baseline   63
3       3             JM  18   M control Baseline   32
4       4             JM  21   F control Baseline   29
5       5             JM  22   F control Baseline   58
6       6             LS  19   F control Baseline   34

```

One more thing to do: an integer has been used for the Subject number, and this variable needs to be turned into a factor before use in a repeated measures design:

```
> longdata$SubF<-as.factor(Subject)
```

This adds a new column to the data where each subject number is now the level of a `SubF` factor.

Repeated Measures ANOVA

With long format data, a simple oneway repeated measures ANOVA can be done using the `aov` command. Unlike the between-subjects designs, the Error term now needs to be added explicitly.

```
> m<-aov(value ~ variable + Error(SubF/variable))
```

The Error term is telling `aov` that each level of `SubF` was measured at each level of `variable`.

```
> summary(m)
```

```

Error: SubF
      Df Sum Sq Mean Sq F value Pr(>F)
Residuals 29  44954    1550

Error: SubF:variable
      Df Sum Sq Mean Sq F value    Pr(>F)
variable  2  27976    13988   13.77 1.28e-05 ***
Residuals 58  58935     1016

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(NB: The `etaSquared` and `TukeyHSD` commands do not work with repeated measures designs.)

“A repeated measures ANOVA showed a significant effect of time $F(2,58)=13.8$, $p<.001$ ”

This isn't very informative as it just shows that the means differ, not whether they go up or down.

If `Subject` had not been turned into the factor `SubF`, this is what happens:

```

> m<-aov(value ~ variable + Error(Subject/variable))
> summary(m)

Error: Subject
      Df Sum Sq Mean Sq F value Pr(>F)
Residuals  1  11721    11721

Error: Subject:variable
      Df Sum Sq Mean Sq
variable  2  27849    13924

Error: Within
      Df Sum Sq Mean Sq F value Pr(>F)
variable  2   2365     1183   1.105  0.336
Residuals 84  89930     1071

```

This non-significant F is because the integer values of `subject` are treated as meaningful values of Subjectness, as if participant 24 were twice as much a Subject than participant 12. **This is clearly nonsense. Please do remember to make sure any Subject variable is being treated as a factor.**

Building a table of means and SDs

To understand what the statistically significant value of F in the oneway repeated measures analysis means, you need some descriptive statistics.

The `aggregate` function computes means and SDs, and its results can be combined into a table, with new column headings, and numbers rounded off sensibly.

```
> tab1<-aggregate(data=longdata, value ~ variable, mean)
```

This makes a table of means of value, split by variable.

```

> tab1
  variable      value
1   Baseline  83.03333
2 After Testing 107.70000
3 Follow-up   126.06667

```

Relabel the value column M to show it holds means

```
> colnames(tab1)[2]="M"
```

The `[2]` makes the command affect the second column of the table.

Add another column called SD by using `aggregate` again to get standard deviations, and asking for the second column of the results by putting `[2]` afterwards.

```

> tab1["SD"]<-aggregate(data=longdata, value ~ variable, sd)[2]
> tab1
  variable      M      SD
1   Baseline  83.03333 34.16944
2 After Testing 107.70000 33.48922
3 Follow-up   126.06667 35.96256

```

Five decimal places is messy, so round columns 2 and 3 of this table to one decimal place:

```

> tab1[,2:3]<-round(tab1[,2:3],digits=1)
> tab1
  variable      M      SD
1   Baseline  83.0  34.2
2 After Testing 107.7 33.5
3 Follow-up   126.1 36.0

```

The `[,2:3]` means the rounding is applied to all rows (the row numbers to operate on are listed before the comma; leaving it blank means all of them) of columns two to three (so `2:5` would affect columns 2 to 5, if they all existed).

These descriptives can help improve the reporting of the repeated measures ANOVA:

“A repeated measures ANOVA showed that scores increased over time $F(2,58)=13.8$, $p<.001$, from $M=83.0$ ($SD=34.2$) at baseline, to $M=107.7$ ($SD=33.5$) at the end of training, rising to $M=126.1$ ($SD=36.0$) at follow-up.”

Mixed measures ANOVA

As well as the within-subject factor, there is a between subject factor.

Again, aov can model this:

```
> m<-aov(value ~ variable * Group + Error(SubF/variable))
> summary(m)
```

Error: SubF

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Group	1	8841	8841	6.854	0.0141 *
Residuals	28	36114	1290		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Error: SubF:variable

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
variable	2	27976	13988	15.100	5.69e-06 ***
variable:Group	2	7058	3529	3.809	0.0281 *
Residuals	56	51877	926		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

“A mixed measures ANOVA with one within-subject factor (time: baseline, after-training, follow-up) and one between subjects factor (condition: control, intervention) showed that overall the intervention group ($M=115.5$, $SD=39.4$) scored higher than the controls ($M=95.7$, $SD=35.3$) $F(1,28)=6.85$, $p=.014$; that scores increased over time (baseline: $M=83.0$, $SD=34.17$; after training: $M=107.7$, $SD=33.5$); follow-up: $M=126.1$, $SD=36.0$) $F(2,56)=15.1$, $p<.001$, and that there was an interaction of time and condition $F(2,56)=3.81$, $p=.028$, with the intervention group’s scores increasing more over time than the control group’s scores (see Figure 1).”

This time the means and SDs came from describeBy commands, part of the psych library – in R there are always lots of ways of doing things. The mat and digits parameters simplify the output by producing a ‘matrix’ instead of listing everything, and only giving two decimal places:

```
> describeBy(value, Group, mat=TRUE, digits=2)
```

item	group1	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
X11	1	control	45	95.69	35.32	94	95.95	40.03	29	164	135	-0.02	-0.94	5.26
X12	2	intervention	45	115.51	39.35	116	115.76	47.44	45	180	135	0.03	-1.15	5.87


```
> describeBy(value, variable, mat=TRUE, digits=2)
```

item	group1	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
X11	1	Baseline	30	83.03	34.17	76.5	81.42	37.81	29	163	134	0.43	-0.60	6.24
X12	2	After Testing	30	107.70	33.49	101.0	107.33	43.00	43	165	122	0.06	-1.06	6.11
X13	3	Follow-up	30	126.07	35.96	123.5	127.38	45.96	46	180	134	-0.20	-0.85	6.57

Plotting an interaction chart

A really quick way to see what is happening in a two-way interaction is this:

```
>interaction.plot(variable,Group,value)
```

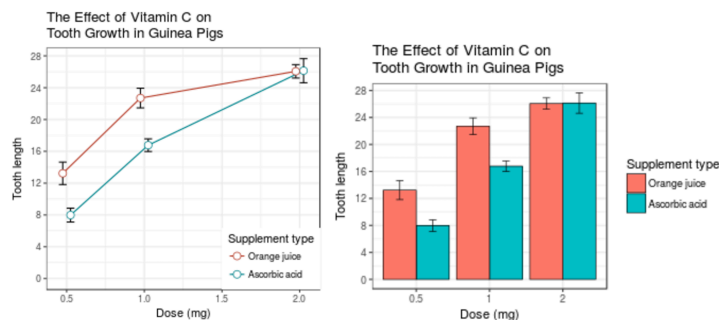


Both groups improve over time, but while the intervention group keep on improving after training, the control group level off. This graph is not good enough for a Results section, though. It also needs to have error bars showing 95% confidence intervals, or standard errors of each mean. These need to be computed and put into a table.

For full details on how to do this, visit :

[http://www.cookbook-r.com/Graphs/Plotting_means_and_error_bars_\(ggplot2\)](http://www.cookbook-r.com/Graphs/Plotting_means_and_error_bars_(ggplot2))

where functions and code can be copied to produce graphs like these:



Using the `longdata` data frame and the `summarySEwithin` functions provided on that page, this command produces the required table of descriptive statistics:

```
> wc.df<-summarySEwithin(longdata,
  measurevar = "value",
  withinvars = "variable",
  betweenvars= "Group",
  idvar=      "Subject",
  na.rm=FALSE, conf.interval = .95)
```

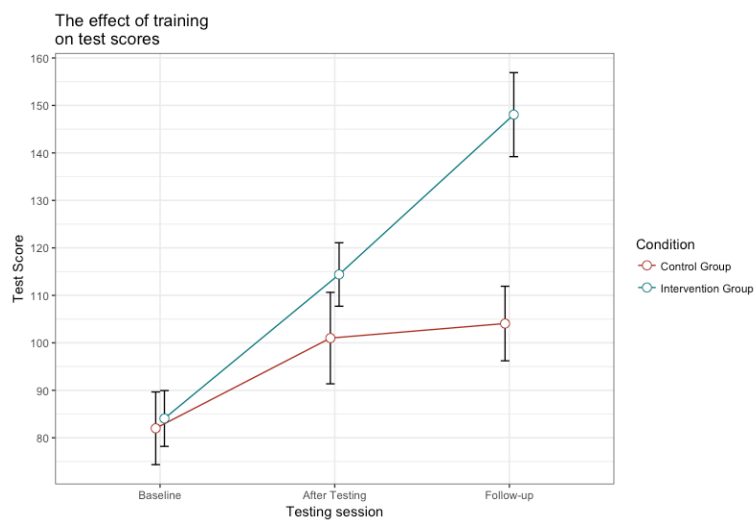
Automatically converting the following non-factors to factors: Group

```
> wc.df
```

	Group	variable	N	value	value_norm	sd	se	ci
1	control	After Testing	15	101.00000	110.91111	37.30264	9.631501	20.65752
2	control	Baseline	15	82.00000	91.91111	29.65254	7.656252	16.42103
3	control	Follow-up	15	104.06667	113.97778	30.39301	7.847441	16.83109
4	intervention	After Testing	15	114.40000	104.48889	25.89478	6.686004	14.34005
5	intervention	Baseline	15	84.06667	74.15556	22.76258	5.877272	12.60550
6	intervention	Follow-up	15	148.06667	138.15556	34.27920	8.850851	18.98319

The following code will then produce an appropriate graph, with standard error bars – the red text shows where variable names or values need to be customised:

```
library(ggplot2)
pd <- position_dodge(0.1) # move series .05 to the left and right to avoid overlap
ggplot(wc.df, aes(x=variable, y=value, colour=Group, group=Group)) +
  geom_errorbar(aes(ymin=value-se, ymax=value+se), colour="black", width=.1, position=pd) +
  geom_line(position=pd) +
  geom_point(position=pd, size=3, shape=21, fill="white") + # 21 is filled circle
  xlab("Testing session") +
  ylab("Test Score") +
  scale_colour_hue(name="Condition", # Legend label, use darker colors
    breaks=c("control", "intervention"),
    labels=c("Control Group", "Intervention Group"),
    l=40) + # Use darker colours, lightness=40
  ggtitle("The effect of training\non test scores") + #\n to force a line-break
  #expand_limits(y=0) + # Uncomment this line to expand y range
  scale_y_continuous(breaks=0:20*10) + # Set y-axis tick every 10 units (max 20 ticks)
  theme_bw() +
  theme(legend.justification=c(1,0),
    legend.position=c(1,0)) # Position legend in bottom right
```



Note that for strict APA format there should be no legend, with the lines being defined in the caption e.g.

Figure 1: Both groups' test scores improved following training, but the intervention group's scores (blue line) continued to improve over the follow-up interval, while the control group's scores (red line) did not. Error bars show one standard error (within subject).

Cronbach's alpha

Exploratory Factor Analysis

Confirmatory Factor Analysis